

# Computer Vision


---

## Class 3 Spectral clustering

- Minimum cut ( $\rightarrow$  Image Processing, class 9):

$$E(u) = \sum_i D(u_i) + \sum_{i,j \in \mathcal{N}(i)} w_{ij} \delta(u_i \neq u_j)$$

- When removing the unary term (the t-links), the optimal solution will assign all nodes to either source or target (no costs)  
 $\rightarrow$  trivial solution



- This is also known as the **shrinking bias** of the minimum cut.
- To avoid the shrinking bias, the volume/connectivity within the resulting region must be considered  
 $\rightarrow$  ratio cut, average cut, normalized cut


- Ratio cut: 
$$\frac{\text{cut}(A, \Omega - A)}{|A|}$$

## Motion segmentation based on point trajectories

- Hard to build a motion model (unary cost) for each object
- Easy to define pairwise cost between point trajectories

## 3D surface reconstruction

- Hard to build a model for the interior of the object volume
- Easy to define costs for potential surface positions (pairwise costs)




- Compute pairwise distances  $d_{ij}$  between all data points (or neighboring data points)
- Greedy heuristic:
  - Merge the two points/clusters with the smallest distance to a single cluster
  - Recompute the affected distances
  - Iterate until remaining distances reach a certain threshold
- Early decisions cannot be corrected  
→ does not optimize a global criterion
- Very sensitive to chosen thresholds






Author: P. Felzenszwalb

- Like agglomerative clustering based on pair-wise distances  $d_{ij}^2$
- Turned into pair-wise affinities  $w_{ij} = \exp(-\lambda d_{ij}^2)$  with a parameter  $\lambda > 0$
- Symmetric  $N \times N$  **affinity matrix**  $W$  ( $N$ : number of data points/pixels)
- Affinity matrix can also be regarded as a fully connected graph with  $N$  nodes
- Groups of points with strong internal connections generate clusters
- Find cut(s) of the graph such that each part is strongly connected (normalized cut vs. minimum cut)



Affinities become edge weights; weakest edges not shown here

- Related segmentation approach for supervised segmentation
- Given a graph and some seed nodes with cluster labels
- At each seed let a particle start and move randomly through the graph.
- Where does it end after infinite time? Repeat the experiment and measure for each node the probability of each particle ending there
- Assign to each node the label with the largest probability (rounding)
- Can be formulated as weighted linear diffusion on the graph (→ efficient computation)



Author: Leo Grady

- Introduce  $K$  labeling vectors  $\mathbf{u}^k$ . Each vector has  $N$  components (= number of nodes in the graph)

- Integer problem with  $u_i^k \in \{0, 1\}$   
Relaxed version with  $u_i^k \in [0, 1]$  (actual random walk)

- Minimize

$$E(\mathbf{u}) = \sum_{k=1}^K \sum_{i=1}^N \sum_{j \in \mathcal{N}(i)} w_{ij} (u_i^k - u_j^k)^2 \quad (\text{gradient descent leads to linear diffusion with diffusivity } w_{ij})$$

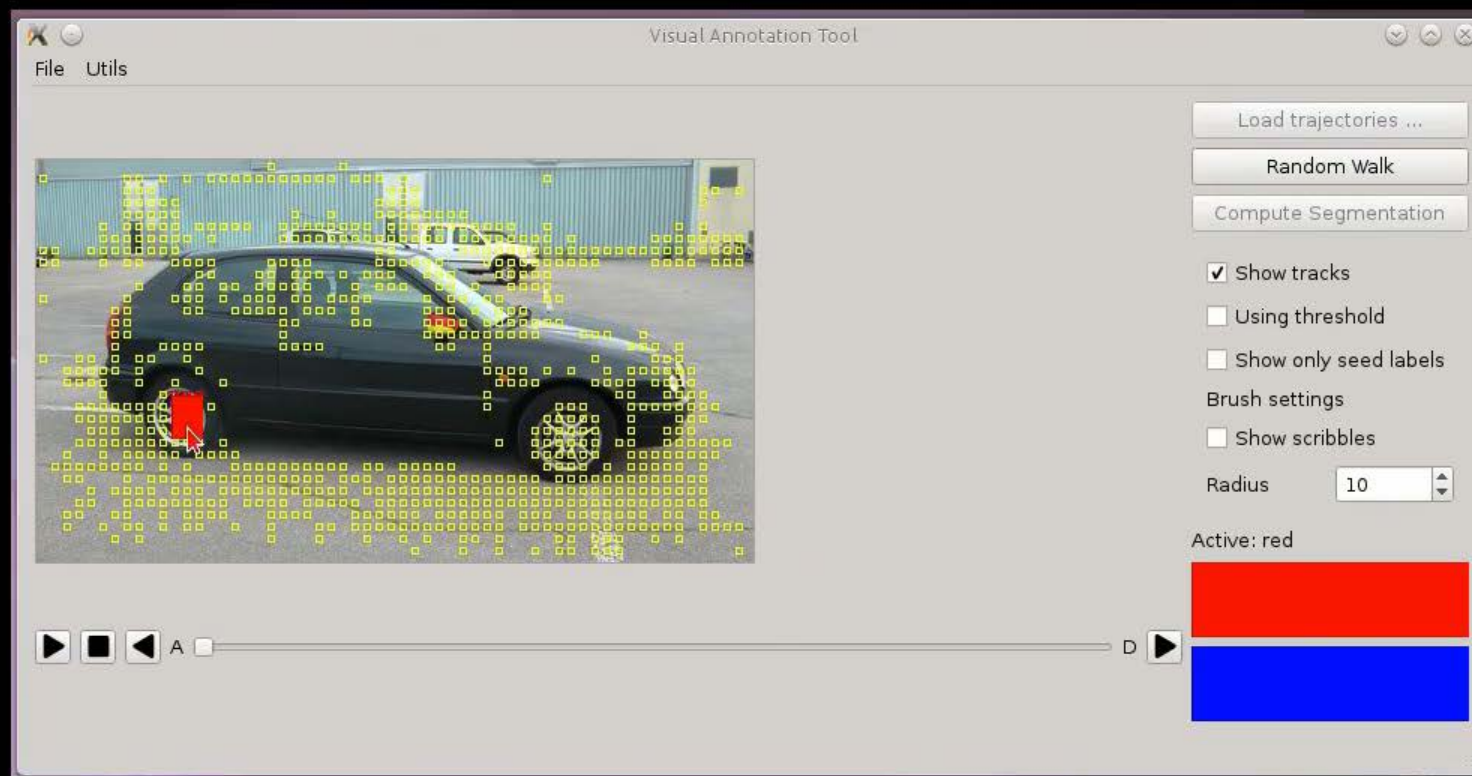
subject to

$$u_i^k = 1 \quad \forall i \in \mathcal{S}^k$$

(keep seed labels  $\mathcal{S}^k$  untouched)

$$\sum_k u_i^k = 1 \quad \forall i$$

(labels must sum to 1)




Video Annotation Tool (Live Demo)

Nagaraja et al. 2015



- Back to the original problem, where we wanted to optimally separate the graph into parts that are themselves well connected
- Two cluster case for the beginning
- Imagine we put a seed at an arbitrary node in the graph
- Running diffusion leads to the equilibrium: all nodes get the seed value  $\rightarrow$  trivial solution
- No surprise: this is the global optimum of
$$E(\mathbf{u}) = \sum_{i=1}^N \sum_{j \in \mathcal{N}(i)} w_{ij} (u_i - u_j)^2$$
- Do not give up too early....



- What is the global optimum that is different from the trivial solution?

$$E(\mathbf{u}) = \sum_{i=1}^N \sum_{j \in \mathcal{N}(i)} w_{ij} (u_i - u_j)^2 \quad \mathbf{u}^\top \mathbf{1} = 0$$


maximally smooth

Orthogonal to the trivial  
solution

- We get a non-trivial solution
- To make this solution independent of the choice of the seed value, enforce a unit length vector:

$$|\mathbf{u}| = 1$$

- Soft separation of the graph
- Can be formulated as an eigenvalue problem



- Define a diagonal matrix  $D$  , where

$$d_i = \sum_j w_{ij}$$

(the sum of all edge weights connected to node  $i$  )

- Graph Laplacian  $D - W$
- Eigenvalue problem on the Laplacian:

$$\mathbf{u}^* = \operatorname{argmin}_{\mathbf{u}} \frac{\mathbf{u}^\top (D - W) \mathbf{u}}{\mathbf{u}^\top \mathbf{u}}$$

- Corresponds to minimizing

$$E(\mathbf{u}) = \sum_{i=1}^N \sum_j w_{ij} (u_i - u_j)^2$$

subject to  $\|\mathbf{u}\| = 1$

- Eigenvalue problem:

$$\mathbf{u}^* = \operatorname{argmin}_{\mathbf{u}} \frac{\mathbf{u}^\top (D - W) \mathbf{u}}{\mathbf{u}^\top \mathbf{u}}$$

- Minimum value is the smallest eigenvalue of  $D - W$
- Corresponding argument is the corresponding eigenvector
- Trivial solution: constant vector with unit length leads to the minimum 0
- More interesting: eigenvector corresponding to the second smallest eigenvalue (by definition orthogonal to the first eigenvector)
- “The smoothest  $\mathbf{u}$  that is orthogonal to the constant vector”

- Computing the eigenvector corresponding to the second smallest eigenvalue of  $D - W$  is a relaxed minimization of the **average cut**

$$\frac{\text{cut}(A, \Omega - A)}{|A|} + \frac{\text{cut}(A, \Omega - A)}{|\Omega - A|}$$

which is a symmetrized version of the **ratio cut**

$$\frac{\text{cut}(A, \Omega - A)}{|A|}$$

- In contrast to the minimum cut, the average cut has a non-trivial global minimum
- This is thanks to the normalization by the size of the emerging regions  $A$  and  $\Omega - A$

- The graph Laplacian is often normalized

$$D - W \quad \rightarrow \quad D^{-\frac{1}{2}}(D - W)D^{-\frac{1}{2}}$$

(weights divided by the geometric mean of the row and column sums)

$$w_{ij} \rightarrow \frac{w_{ij}}{\sqrt{d_i d_j}}$$

- The eigenvector  $\mathbf{z}$  of the normalized problem must be rescaled:

$$\mathbf{u} = D^{-\frac{1}{2}}\mathbf{z} \quad (\text{see Shi-Malik 2000 for details})$$

- This leads to the **normalized cut**



$$\frac{\text{cut}(A, \Omega - A)}{\text{assoc}(A, \Omega)} + \frac{\text{cut}(A, \Omega - A)}{\text{assoc}(\Omega - A, \Omega)}$$

Measures the removed edges (cut) relative to the total edge weight originating in each region (assoc)

- Minimizing the normalized cut is equivalent to maximizing the **normalized association**

$$\frac{\text{assoc}(A, A)}{\text{assoc}(A, \Omega)} + \frac{\text{assoc}(\Omega - A, \Omega - A)}{\text{assoc}(\Omega - A, \Omega)}$$


- Advantage I: Limiting case where looking for similarities is the same as looking for differences
- Advantage II: Can separate areas with higher connectivity from areas with lower connectivity





- In many applications, the affinity matrix is large  
→ usual numerical methods for computing eigenvalues with cubic complexity  $\mathcal{O}(N^3)$  are prohibitively slow
- 1. Some graphs may not be fully connected  
→ matrix will be sparse (many zero entries)
- 2. We just require the smallest eigenvalues and their corresponding eigenvectors, not all of them
- Can we exploit this to have a faster solver?
- Indeed there is the **Lanczos method**, which efficiently computes just the smallest (or largest) eigenvalues of a matrix in  $\mathcal{O}(N^2)$
- Available via the ARPACK numerical library (wrapped by SciPy via “eigs”)



- Recall: We solve the real-valued relaxed problem (spectral relaxation)  
→ generally no clear cut of the graph, just soft indicators
- We get an approximate solution of the binary segmentation problem (average cut or normalized cut) by thresholding the eigenvector (rounding)




Input image

2<sup>nd</sup> Eigenvector

After thresholding

- Average cut and normalized cut are NP hard problems

- The concept can be extended to more than two clusters
- The eigenvector to the third smallest eigenvalue indicates an alternative partitioning of the graph
- In the general case we compute the  $K$  smallest eigenvalues
- The corresponding eigenvectors span a  $K$ -dimensional subspace
  - Each node (data point) maps to a  $K$ -dimensional vector
  - Mapping is called **Laplacian eigenmap**
- We can run standard clustering techniques (k-means) to convert the real-valued vectors into integer labels  
→ **spectral clustering**




Mappings from a graph to its eigenspace, clustering in the eigenspace


- The balancing property of the average and normalized cut transfers to spectral clustering, too.

$$\frac{\text{cut}(A, \Omega - A)}{|A|} + \frac{\text{cut}(A, \Omega - A)}{|\Omega - A|} \quad (\text{average cut})$$

- Consequence: When affinities are not informative, spectral clustering tends to create equally sized regions
- With constant affinities, the eigenvectors corresponding to the smallest eigenvalues yield the Fourier basis
- Consider balancing when applying spectral clustering on non-uniform grids



Eigenvectors of constant affinity matrix



After thresholding

- Another graph based technique is to find the most connected subset among the nodes:

$$\max_A \left( \frac{\text{assoc}(A, A)}{|A|} \right)$$

- Can be used to find outliers in a dataset
- Relaxed problem: find the largest eigenvector of the affinity matrix  $W$
- Nodes of the most connected subset obtain non-zero values
- Largest eigenvector can be found also by **power iteration** if  $\lambda_1 > \lambda_2$   
$$v^{k+1} = \frac{Wv^k}{\|Wv^k\|} \text{ (very simple algorithm)}$$
- Major part of the PageRank algorithm that initiated Google's success

- L. Grady: Random walks for image segmentation, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 28(11):1768-1783, 2006.
- J. Shi, J. Malik: Normalized cuts and image segmentation, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 22(8):888-905, 2000.
- F. R. K. Chung, *Spectral Graph Theory*, American Mathematical Society, 1997.
- A. Ng, M. Jordan, Y. Weiss: On spectral clustering: analysis and an algorithm, *Advances in Neural Information Processing Systems (NIPS)*, 2001.
- P. Perona, W. Freeman: A factorization approach to grouping, *European Conference on Computer Vision*, pp. 655-670, 1998.